

Exam 2 review

Monday, April 8, 2024 1:00 PM

Chapter 3 Sections 3.1 - 3.6

Addition and subtraction will be the same thing. Be careful with overflow and underflow!

Multiplication

Start with long-multiplication approach

The normal approach: result is the sum of operand lengths

Binary multiplication is: bunch of shifts and adds

Note that: n-bit number multiply with n-bit number will result in 2n-bit number

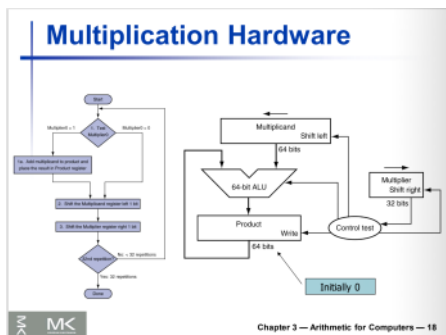
Integer multiplication

Key elements

1. Examine multiplier bits from right to left
2. Shift multiplicand left one position each step
3. Simplification: each step, add multiplicand to running product total, but only if multiplier bit == 1

Detail step:

1. Initialize product register to 0
2. Get result of the multiplier bit then add it to product (Ignore if bit is 0)
3. Shift multiplicand left
4. Move to the next one



Can optimize this by parallel operation add and/shift

LEGv8 multiplication

Three multiply instructions:

- MUL: multiply
 - Gives the lower 64 bits of the product
- SMULH: Sign multiply high
 - Gives the upper 64 bits of the product, assuming the operands are signed
- UMLH: unsigned multiply high
 - Gives the lower 64 bits of the product, assuming the operands are unsigned

Floating point

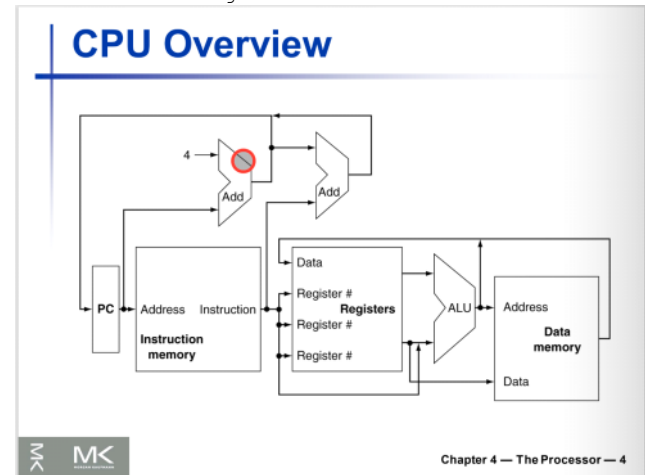
Follow normal instruction and rmb some

Chapter 4

Read all except 4.6 to 4.9 and 4.12

Sequence for Instruction Execution


- PC -> Instruction memory, fetch instruction
- Register numbers -> register file, read registers
- Depending on instruction class
 - Use ALU to calculate
 - Arithmetic result
 - Memory address for load/store
 - Branch target address
 - Access data memory for load/store
 - PC <- target address or PC + 4



- To join wires, use multiplexers (MUX)
- Add control to control the sequence of activation for every components

Clocking Methodology

- Combinational logic transforms data during clock cycles
 - Between clock edges
 - Input from state elements, output to state element
 - Longest delay determines clock period



Data-path

Elements that process data and addresses in the CPU

R-format instructions

- Read two register operands
- Perform arithmetic/logical operation
- Write register result

Using: Register file (R1, R2) and ALU

Load/store instructions

- Read register operand
- Calculate address using 16-bit offset
 - Use ALU, but sign-extend offset
- Load: read memory and update register
- Store: write register value to memory

Branch instructions

- Read register operands
- Compare operands
 - Use ALU, subtract and check Zero output
- Calculate target address
 - Sign-extend displacement

are unsigned

Floating point

Follow normal instruction and rmb some special case such as infinity, overflow and underflow

IEEE floating-point format

Bias table

	Single
Exponent	8
Fraction	23
Bias	127

Rmb how to normalize and denormalize numbers as well

Denormalize usually work for underflow number

Special cases:

A	Exp	Fraction
Infinities	All 1	All 0
NaNs	All 1	Not all 0

Learn again about how to add two floating-point

?Rounding with guard bit

Subword parallelism

Graphics and audio application can take advantage of performing simultaneous operations on short vectors

Example: 128-bit adder can perform in
Sixteen 8-bit adds
Eight 16-bits adds
Four 32-bit adds

Also called data-level parallelism, vector parallelism, or Single Instruction, Multiple Data (SIMD)

Fallacies and pitfalls

Generally stem from the difference between the limited precision of computer arithmetic and the unlimited precision of natural arithmetic.

Fallacy: Just as a left shift instruction can replace an integer multiply by a power of 2, a right shift is the same as an integer division by a power of 2

=> Shift -> mul/div by 2^x => But only work with unsigned num

Pitfall: Floating-point addition is not associative

When adding two large number of opposite signs plus a small number
 $C = -1.5 * 10^{30}$
 $A = 1.5 * 10^{30}$
 $B = 1$

$C + (A + B) = -1.5 * 10^{30} + (1.5 * 10^{30} + 1.0) = 0$

$(C + A) + B = (-1.5 * 10^{30} + 1.5 * 10^{30}) + 1.0 = 1$

NOPE THEY ARE NOT THE SAME. EVEN THO

- Compare operands
 - Use ALU, subtract and check Zero output
- Calculate target address
 - Sign-extend displacement
 - Shift left 2 places (word displacement)
 - Add to PC + 4
 - Already calculated by instruction fetch

Pipelining

Literally, overlapping execution.

Pipelining will get the longest time for a component to process the data and use it as the time for all other processes.

LEGv8 Pipeline

1. IF: Instruction fetch from memory
2. ID: Instruction decode & register read
3. EX: execute operation or calculate address
4. MEM: access memory operand
5. WB: Write result back to register

4.10. Parallelism via Instructions

NOTICE: complex topics

?????

4.14. Fallacies and pitfalls

Fallacy: pipelining is easy

Fallacy: Pipelining ideas can be implemented independent of technology

-> Today, power are leading to less aggressive and more efficient designs

Pitfall: failure to consider instruction set design can adversely impact pipelining

4.15. Concluding remarks

THEY SHOULD BE

*Fallacy: parallel execution strategies
that work for integer data types also
work for floating-point data types*

*Fallacy: only theoretical
mathematician care about floating-
point accuracy*

3.11 Concluding remarks

- Bits have no inherent meaning
 - Interpretation depends on the instructions applied
- Computer representations of numbers
 - Finite range and precision
 - Need to account for this in programs
- ISAs support arithmetic
 - Signed and unsigned integers
 - Floating-point approximation to reals
- Bounded range and precision
 - Operations can overflow and underflow